

# Data Processing, Analysis and Visualization in Python

Basic Machine Learning II –  
Unsupervised Learning



1

## Outline

- Introduction
- Dimensionality Reduction, Feature Extraction, and Manifold Learning
  - Principal Component Analysis (PCA)
  - Manifold Learning with t-SNE (t-distributed Stochastic Neighbor Embedding)
  - UMAP (Uniform Manifold Approximation and Projection)
- Clustering
  - K-Means Clustering



2

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



3

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
  - Isomap
  - Locally Linear Embedding
  - Modified Locally Linear Embedding
  - Hessian Eigenmapping
  - Spectral Embedding
  - Local Tangent Space Alignment
  - Multi-dimensional Scaling (MDS)
  - t-distributed Stochastic Neighbor Embedding (t-SNE)
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



4

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
  - K-means
  - Affinity Propagation
  - Mean Shift
  - Spectral clustering
  - Hierarchical clustering
  - DBSCAN
  - OPTICS
  - BIRCH
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



5

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
  - Spectral Co-Clustering
  - Spectral Biclustering
  - Biclustering evaluation
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



6

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
  - Principal component analysis (PCA)
  - Kernel Principal Component Analysis (kPCA)
  - Truncated singular value decomposition and latent semantic analysis
  - Dictionary Learning
  - Factor Analysis
  - Independent component analysis (ICA)
  - Non-negative matrix factorization (NMF or NNMF)
  - Latent Dirichlet Allocation (LDA)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



7

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
  - Empirical covariance
  - Shrunk Covariance
  - Sparse inverse covariance
  - Robust Covariance Estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)



8

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
  - Overview of outlier detection methods
  - Novelty Detection
  - Outlier Detection
  - Novelty detection with Local Outlier Factor
- Density estimation
- Neural network models (unsupervised)



9

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
  - Density Estimation: Histograms
  - Kernel Density Estimation
- Neural network models (unsupervised)



10

## Unsupervised Learning in Scikit-learn

- Gaussian mixture models
- Manifold learning
- Clustering
- Biclustering
- Decomposing signals in components (matrix factorization problems)
- Covariance estimation
- Novelty and outlier detection
- Density estimation
- Neural network models (unsupervised)
  - Restricted Boltzmann machines



11

## Principal Component Analysis (PCA)


- The central idea of principal component analysis (PCA) is to reduce the dimensionality of a data set consisting of a large number of interrelated variables while retaining as much as possible of the variation present in the data set. This is achieved by transforming to a new set of variables, the *principal components (PCs)*, which are uncorrelated (i.e., orthogonal), and which are ordered so that the first few retain most of the variation present in all of the original variables.
- PCA can be thought of as an **unsupervised learning** problem.

12

## PCA in a nutshell

- Let's assume we have 6 datapoints and 2 variables


	residual sugar	alcohol
1	2.6	9.8
2	2.3	9.8
3	1.9	9.8
4	1.9	9.4
5	1.8	9.4
6	1.6	9.4



13

## PCA in a nutshell


- Calculate the mean of each variable
- Calculate the center of the data for each variable and shift to origin



14

## PCA in a nutshell


- Draw a line through the origin and then rotate until it fits the points optimally



15

## PCA in a nutshell


- Draw a line through the origin and then rotate until it fits the points optimally. How do we determine the optimal rotation?



16

## PCA in a nutshell


Draw a line through the origin and then rotate until it fits the points optimally. This is PC<sub>1</sub>.



17

## PCA in a nutshell


1. Draw a line through the origin and then rotate until it fits the points optimally. This is PC<sub>1</sub>.



18

## PCA in a nutshell


The line orthogonal to PC<sub>1</sub> through the origin is PC<sub>2</sub>.



19

## PCA in a nutshell

To get the projected points, we rotate back the PC axes to original coordinate set



20

## PCA in a nutshell

How important is each PC?

$SS_1$  (distances for  $PC_1$ ) = eigenvalue for  $PC_1$

$SS_2$  (distances for  $PC_2$ ) = eigenvalue for  $PC_2$

Convert the eigenvalues to variation for each PC:

$$\frac{SS_1}{N-1} = \text{Variance ratio } PC_1 = \sigma_1 \Rightarrow \frac{\sigma_1}{\sigma_1 + \sigma_2}$$

$$\frac{SS_2}{N-1} = \text{Variance ratio } PC_2 = \sigma_2 \Rightarrow \frac{\sigma_2}{\sigma_1 + \sigma_2}$$

How much does each PC contribute to the variation in the data

21

## The mathematics of PCA

- Compute the *covariance matrix*  $\text{cov}(X, X) = C$  of the whole dataset.

$$C = \frac{\sum_{i=1}^N (x_i - \bar{x}) \sum_{i=1}^N (x_i - \bar{x})^T}{N-1}$$

n\_samples  
vector of size n\_features

- Solve eigen-equation

$$Cv = \lambda v$$

Covariance matrix (n\_features × n\_features)  
Eigenvector of C (n\_features × n\_features)  
Eigenvalue of v (n\_features)

- by solving determinant determine eigenvalues

$$\det(C - \lambda I)$$

The eigenvalues are the roots of C

- Determine eigenvectors by substituting v into the eigen-equation
- Form a matrix W with the k highest eigenvectors and use W as follows:

$$x' = W^T \cdot x$$

Transformed samples onto new subspace

22

## Principal Component Analysis (PCA)

- The process of obtaining principal components from a raw dataset can be simplified in 5 parts :

1. Take the whole dataset consisting of **n\_features+1 dimensions** and ignore the labels such that our new dataset becomes **n\_features dimensional**. Compute the *mean* for every dimension of the whole dataset.
2. Compute the *covariance matrix* of the whole dataset.
3. Compute *eigenvalues* and the corresponding *eigenvectors*.
4. Sort the eigenvectors by decreasing eigenvalues and choose **k** eigenvectors with the largest eigenvalues to form a **n\_features × k dimensional matrix W**.
5. Use this **n\_features × k eigenvector matrix** to transform the samples onto the new subspace.

23

## PCA NumPy version

```
import numpy as np
import numpy.linalg as linalg
from sklearn.datasets import make_classification

n_features = 2
X = np.array([[1, 2], [3, 4], [5, 6]])
n_samples = X.shape[0]
# We center the data and compute the sample covariance matrix
X_centered = (X - np.mean(X, axis=0))
# Compute covariance matrix
cov_matrix = np.dot(X_centered.T, X_centered) / (n_samples - 1)
# Eigendecomposition of covariance matrix
eigenvalues, eigenvectors = linalg.eig(cov_matrix)
# Sort eigenvalues and associated eigenvectors using index-based sorting
idx = eigenvalues.argsort()[::-1]
eigenvalues = eigenvalues[idx]
eigenvectors = eigenvectors[:,idx]
# Eigenvectors corresponding to the k maximum eigenvalues
W = eigenvectors[:,0:2]
# Transform the samples onto the new subspace
X_transformed = np.dot(W.T, X_centered.T)
```

pca\_numpy.ipynb

24

## PCA NumPy version

```
print("Original dataset\n", X)

Original dataset
[[1 2]
 [3 4]
 [5 6]]

print("Transformed dataset\n", X_transformed.T)

Transformed dataset
[[-2.82842712  0.          ]
 [ 0.          0.          ]
 [ 2.82842712  0.          ]]
```

25

## PCA sklearn version

```
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

pca = PCA()
pca.fit(X)

# access values and vectors
eigenvalues_pca = pca.explained_variance_
eigenvectors_pca = pca.components_
X_transformed = pca.fit_transform(X)
print("Original dataset\n", X)
Original dataset
[[1 2]
 [3 4]
 [5 6]]

print("Transformed dataset\n", X_transformed)

Transformed dataset
[[-2.82842712  0.          ]
 [ 0.          0.          ]
 [ 2.82842712  0.          ]]
```

pca\_numpy.ipynb

26

## Principal Component Analysis (PCA)

```
>>> import pandas as pd

>>> df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-
learning-databases/wine-quality/winequality-red.csv',
delimeter=';')
>>> df['hue'] = 'red'
>>> df2 = pd.read_csv('https://archive.ics.uci.edu/ml/machine-
learning-databases/wine-quality/winequality-white.csv',
delimeter=';')
>>> df2['hue'] = 'white'
>>> df_wine = pd.concat([df, df2], ignore_index=True)
>>> df_wine
```

pca\_kmeans\_wine\_color.ipynb

27

## Principal Component Analysis (PCA)

	n_features →											
	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5 red
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5 red
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5 red
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6 red
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5 red
...	...	...	...	...	...	...	...	...	...	...	...	...
6492	6.2	0.21	0.29	1.6	0.039	24.0	92.0	0.99114	3.27	0.50	11.2	6 white
6493	6.6	0.32	0.36	8.0	0.047	57.0	168.0	0.99490	3.15	0.46	9.6	5 white
6494	6.5	0.24	0.19	1.2	0.041	30.0	111.0	0.99254	2.99	0.46	9.4	6 white
6495	5.5	0.29	0.30	1.1	0.022	20.0	110.0	0.98869	3.34	0.38	12.6	7 white
6496	6.0	0.21	0.38	0.8	0.020	22.0	98.0	0.98941	3.26	0.32	11.8	6 white

n\_samples ↓

6497 rows × 13 columns

Principal components are new variables that are constructed as linear combinations or mixtures of the initial variables.

$$PC_k = \sum_{i=1}^{n_{features}} v_i f_i \quad \forall k$$


$v_i$  ← eigenvector  $i$  (i.e., coefficients)  
 $f_i$  ← feature  $i$

28

## Principal Component Analysis (PCA)



```
>>> import seaborn as sns
>>> ax = sns.boxplot(data=df_wine, orient='h', palette='PuBuGn')
```




29

## Principal Component Analysis (PCA) Data standardization



```
>>> from sklearn.preprocessing import StandardScaler
>>> data_pca = df_wine.copy()
>>> data_pca = data_pca.drop(labels = ['quality', 'hue'], axis = 1)
>>> # StandardScaler() scales the data. The fit_transform() module fits these new
>>> # values to the data, and stores them, replacing the old values.
>>> data_pca = StandardScaler().fit_transform(data_pca)
>>> # Plot the transformed (scaled and centered) data:
>>> ax = sns.boxplot(data=data_pca, orient='h', palette='PuBuGn')
```



30

## Principal Component Analysis (PCA)




```
>>> from sklearn.decomposition import PCA
>>> # Apply PCA on the transformed (scaled and centered) data:
>>> pca = PCA(n_components=2)
>>> print(type(pca))
<class 'sklearn.decomposition._pca.PCA'>
>>> X = pca.fit_transform(data_pca)
>>> X.shape
(6497, 2)
>>> X
array([[ -3.20599617,  0.41652332],
       [ -3.03905081,  1.10746213],
       [ -3.07189347,  0.87896444],
       ...,
       [ 0.5711325 , -0.72266165],
       [ 0.09005243, -3.54577991],
       [ 0.51257566, -2.89104008]])
```

31

## Principal Component Analysis (PCA)



```
>>> import matplotlib.pyplot as plt
>>> pca_dataset = pd.DataFrame(data = X, columns = ['component1', 'component2'])
>>> pca_dataset['hue'] = df_wine['hue']
>>> plt.figure()
>>> plt.figure(figsize=(6,6))
>>> plt.xlabel('Component 1')
>>> plt.ylabel('Component 2')
>>> plt.title('2 Component PCA')
>>> sns.scatterplot(x = pca_dataset['component1'], y = pca_dataset['component2'],
                  hue=pca_dataset['hue'], alpha=0.3, palette='PuBuGn')
```



Similar datapoints are closer together, forming a cluster. For this dataset we see that white wine and red wine form two separate clusters.

32



## Principal Component Analysis (PCA)

```
>>> # Trying to decipher the meaning of the principal components
>>> print("Meaning of the 2 components:")
>>> eigenvectors = pca.components_ # The eigenvectors
>>> print(eigenvectors.shape)
(2, 11)
>>> pd.DataFrame(eigenvectors,
                  columns=df_wine.keys().drop(labels = ['quality', 'hue']))
```

	n_features										
n_components	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	-0.239850	-0.378601	0.147806	0.344812	-0.289520	0.433571	0.488749	-0.045997	-0.215840	-0.295267	-0.108741
1	0.335908	0.120606	0.183417	0.334376	0.311061	0.068348	0.086109	0.583883	-0.160638	0.190101	-0.464181


```
>>> eigenvalues = pca.explained_variance_
>>> print(pd.Series(eigenvalues))
```

```
0    3.043073
1    2.494307
dtype: float64
```

33

## Principal Component Analysis (PCA)


```
>>> # How much does each component explain the original dataset.
>>> string = [s+1 for s in range(len(pca.explained_variance_ratio_))]
>>> variance_ratio = pd.DataFrame(pca.explained_variance_ratio_,
                                  columns=['variance_ratio'], index=string)
>>> ax = sns.barplot(data=variance_ratio, y='variance_ratio',
                    x=variance_ratio.index, palette='PuBuGn')
>>> ax.set(xlabel='Principal components', ylabel='Variance ratio')
```



Try it yourself!  
Try with different number of PCs

34


## PCA in the Cu-binding protein CueR



35

## K-means Clustering

How do we cluster points in space?



36

## K-means Clustering

How do we cluster points in space?



37

## K-means Clustering

How do we cluster points in space?



38


## K-means Clustering

Step 1: Select the number of clusters. Let's try 3.

Step 2: Select 3 random clusters

Step 3: Measure distance between all points and all clusters

Step 4: Assign each point to the nearest cluster



39

## K-means Clustering

In this case we get the following cluster:




Step 5: Calculate the centroid of each cluster

Step 6: Calculate the distance of all points relative to the centroids and re-cluster. If assignment of points doesn't change, we're done.

40

## K-means Clustering

Step 7: Now we calculate the variation within each cluster:



Total variation within the clusters

This clearly isn't an optimal cluster, but K-means doesn't know this. Solution?

41


## K-means Clustering

Step 1: Select a **new** set of clusters. Let's try 3.

Step 2: Select 3 random clusters

Step 3: Measure distance between all points and all clusters


Step 4: Assign each point to the nearest cluster



42

## K-means Clustering

In this case we get the following cluster:



Step 5: Calculate the centroid of each cluster

Step 6: Calculate the distance of all points relative to the centroids and re-cluster. If assignment of points doesn't change, we're done.

43

## K-means Clustering

In this case we get the following cluster:



Step 5: Calculate the centroid of each cluster


Step 6: Calculate the distance of all points relative to the centroids and re-cluster. If assignment of points doesn't change, we're done.

44

## K-means Clustering



Step 7: Now we calculate the variation within each cluster:




45

## K-means Clustering



Since K-means doesn't know which is the best solution, it will do a certain number of clusters and determine which is the optimal (best distribution of variation within clusters). Let's say we tried 3 random initial guesses:



46

## K-means Clustering



What's the optimal value for K?

Start with K=1 and then increase the number of K's.

K=1:



47


## K-means Clustering



What's the optimal value for K?

Start with K=1 and then increase the number of K's.

K=2:



48


## K-means Clustering



What's the optimal value for K?

Start with K=1 and then increase the number of K's.

K=3:



49

## K-means Clustering



What's the optimal value for K?

Start with K=1 and then increase the number of K's.

K=4:




50

## K-means Clustering



What's the optimal value for K?

Comparing the variation with K:



51

## K-means Clustering



What's the optimal value for K?

Plotting the variation vs. K (elbow plot):




We'll see in a bit that we'll use something called inertia instead of variation

52

## K-means Clustering

How do we cluster points in 2D, 3D, ...,  $ND$  space?  
Same idea!



53

## K-means Clustering

We have to decide how to compute the distance between points.

The K-means algorithm aims to choose centroids that minimize the **inertia**, or **within-cluster sum-of-squares** criterion:


$$\sum_{i=0}^N \min_{\mu_j \in C} (\|x_i - \mu_j\|^2)$$

All points in cluster  $j$  (pointing to  $x_i$ )  
 Centroid  $j$  (pointing to  $\mu_j$ )  
 Sum over points in cluster,  $j$  (pointing to the summation symbol)

54

## K-means Clustering

1. Guess number of clusters  $K$
2. Form the initial clusters
3. Iterate until convergence





K-Means  
in action

Source: <http://mcla.ue/blog/k-means-clustering.html>


55

## Problematic cases for K-means


Incorrect number of blobs    Anisotropically distributed blobs



Unequal variance

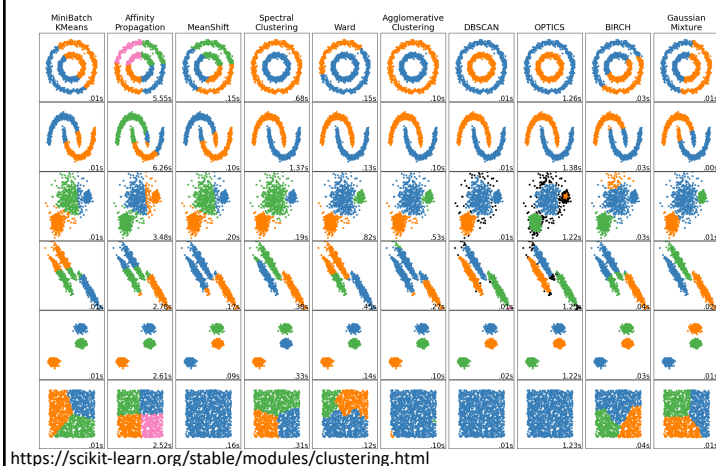


Unevenly sized blobs



56

## Behavior of clustering methods



57


## K-means example

```
from sklearn.cluster import KMeans
inertia = []
# Creating 10 K-Mean models while varying the number of clusters (k)
# An elbow in the graph indicates the right number of clusters
for k in range(1,11):
    model = KMeans(n_clusters=k, init='k-means++')
    model.fit(pca_dataset.iloc[:,2]) # Fit model to samples
    inertia.append(model.inertia_) # Append the inertia to
                                   # the list of inertias
inertia = pd.DataFrame({'Inertia':inertia}, index=range(1,11))
ax = sns.lineplot(data=inertia, marker='o', palette='PuBuGn')
ax.set(xlabel='Number of clusters, k', ylabel='Inertia')
```

pca\_kmeans\_wine\_color.ipynb

58


## K-means example



59

## K-means example

```
model = KMeans(n_clusters=3, init='k-means++')
model.fit(pca_dataset.iloc[:,2])
labels = model.predict(pca_dataset.iloc[:,2])
sns.scatterplot(x = pca_dataset[0], y = pca_dataset[1],
                alpha=0.3, hue=labels, palette='Set1')
ax.set(xlabel='PC1', ylabel='PC2')
```





60

## K-means example



```
centers = model.cluster_centers_ # Get the cluster centroids  
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=200,  
            alpha=0.5) # Plot centroids on top of clusters
```



 Try it yourself!  
Try with different number of clusters

61